

Homework 3

Carl Eastlund

Due **Wed., Feb. 27 at 9:00pm.**

1. Associative Maps.

(a) An AssocMap is a Tree:

```
;; A Tree is one of:  
;; - empty  
;; - (node Integer Key Value Tree Tree)  
(struct node [height key value left right] #:transparent)  
;; where:  
;; - for any node N in left, left.key < key  
;; - for any node N in right, key < right.key  
;; - height = 1 + max(left.height, right.height)  
;; - left.height <= right.height + 1  
;; - right.height <= left.height + 1  
  
;; A Key is a Number  
;; A Value is a String
```

Note: Because the left and right subtrees of a node must have heights within 1 of each other, the maximum height of a tree is $O(\log n)$. The argument is based on the fact that we can construct the smallest possible tree for any height n . For height 0, the tree is empty. For height 1, the tree is a singleton node. For any higher height n , the smallest possible tree is a node with the smallest tree of height $n - 1$ on the left and the smallest tree of height $n - 1$ on the right. Adding the solutions for $n - 1$ and $n - 2$ suggest the Fibonacci numbers; we can verify that the smallest tree of height n has a number of leaves equal to one less than the $n + 1$ st Fibonacci number. Since the Fibonacci numbers grow exponentially, the minimum number of nodes for a given height therefore grows exponentially. The maximum height for a given number of nodes is the inverse of this function, and it therefore grows logarithmically. Interestingly, we get a logarithmic *height* bound without guaranteeing any constant factor bound for the relative *size* of the left and right subtrees.

- (b)
 - i. Every key/value pair occupies a node, which takes $\Theta(1)$ space. Because keys are strictly increasing from left to right, there can be no duplicate keys, so there are n nodes for n keys. Therefore the total space is $\Theta(n)$.
 - ii. `fresh-assoc` uses `empty-tree` which produces empty, which is trivially valid.
 - iii. `assign` uses `tree-insert`. The `tree-insert` function recurs left or right depending on the given key, which preserves the order invariants on the tree. It builds nodes with `balanced-node` and `almost-balanced-node`, which automatically construct the height field correctly. Insertion can increase the height of a tree by at most one; `almost-balanced-node` preserves the balance invariants for trees that are off-balance by at most one by performing one or two tree rotations as needed. Therefore `tree-insert` produces a valid Tree at every step.
 - iv. `unassign` uses `tree-delete-range` for a singleton range. The function `tree-delete-range` uses `tree-append`, `tree-filter<`, and `tree-filter>`. These functions all preserve the key order from the original tree. They all ultimately use `unbalanced-node`, `almost-balanced-node`,

or balanced-node to construct nodes, which all maintain the height invariants. The unbalanced-node function recurs on the left or right to “push” the smaller subtree down inside the larger one to find siblings of similar size, then calls almost-balanced-node when it gets there. The result is always a balanced tree.

- (c)
 - i. We use worst-case analysis for our associative mappings.
 - ii. fresh-assoc produces empty immediately.
 - iii. assign calls tree-insert, which recurs $O(\log n)$ times and only calls $\Theta(1)$ -time functions.
 - iv. unassign calls tree-append, tree-filter<, and tree-filter>. These functions all recur at most $O(\log n)$ times. They all also call unbalanced-node. The running time of unbalanced-node is odd; it runs in $O(|\log m - \log n|)$ time for trees of size m and n . Its running time is proportional to the *difference* between the heights of its arguments. As tree-append, tree-filter<, and tree-filter> recur, their calls to unbalanced-node “walk” down the height of a tree. If one call to unbalanced-node takes several steps by going from height $x+k$ to height x , the next call will start at height x . In other words, the *total* time spend in calls to unbalanced-node is $O(\log n)$ for any of tree-append, tree-filter<, and tree-filter>. Therefore all three functions take $O(\log n)$ time.
 - v. lookup calls tree-search which recurs at $O(\log n)$ times and calls no helpers.
- 2. (a) A Set is a Tree as defined above, in which every key represents an element of the set and the values associated with keys are irrelevant.
- (b)
 - i. Trees take $\Theta(n)$ space as argued above.
 - ii. empty-set calls empty-tree, which produces a valid tree as argued above.
 - iii. extend calls tree-insert, which produces a valid tree as argued above.
 - iv. without calls tree-delete-range, which produces a valid tree as argued above.
- (c)
 - i. Our bounds are all worst-case.
 - ii. empty-set takes $O(1)$ time trivially.
 - iii. in? calls tree-search, which takes $O(\log n)$ time as argued above.
 - iv. extend calls tree-insert, which takes $O(\log n)$ time as argued above.
 - v. without calls tree-delete-range, which takes $O(\log n)$ time as argued above.
- 3. (a) ;; A Queue is (queue Chain Chain).
 (struct queue [leftmost rightmost] #:mutable #:transparent)
 ;; Either leftmost and rightmost are both empty, or they are both links.
 ;; If both are links, then they are part of a single chain of links:
 ;; * leftmost.left = empty
 ;; * rightmost.right = empty
 ;; * for any link L other than leftmost, L.left.right = L
 ;; * for any link L other than rightmost, L.right.left = L
 ;; * from any link L, leftmost is reachable by going left 0 or more times
 ;; * from any link L, rightmost is reachable by going right 0 or more times
 ;; These invariants must hold before and after any Queue operations.
 ;; Note that during Queue operations, in between mutations,
 ;; some or all of these invariants may be violated.

 ;; A Chain is either empty or (link Link String Link).
 (struct link [left elem right] #:mutable #:transparent)
- (b)
 - i. A queue contains one link per value, and every link takes $\Theta(1)$ space. The total is therefore $\Theta(n)$ space.
 - ii. new-queue produces a queue that is trivially valid.
 - iii. add-leftmost maintains the left and right link invariants when adding links, and therefore produces a valid queue.

- iv. `drop-leftmost` maintains the left and right link invariants when removing links and correctly detects when the queue becomes empty; it therefore produces a valid queue.
 - v. `add-rightmost` maintains the left and right link invariants when adding links, and therefore produces a valid queue.
 - vi. `drop-rightmost` maintains the left and right link invariants when removing links and correctly detects when the queue becomes empty; it therefore produces a valid queue.
 - vii. `concatenate` maintains the left and right link invariants, and ensures that every link is part of only one queue; it therefore produces a valid queue.
- (c) The analysis for queues is worst-case, and by inspection every operation takes $\Theta(1)$ time because there is no recursion or iteration.